

# Generative Adversarial Network based heuristics for sampling based path planning Project description\*

Azamat Kanametov

Alina Kolesnikova

Timofey Zinenko

May 11, 2021

## 1 Optimal path planning problem

**Non-formal** Determine a collision-free path from a start point to a goal point while optimizing a performance criterion such as distance, time or energy

**Formal**

**Given:**

- $\mathcal{X} \in \mathbb{R}^n$  – the state space,  $n \in \mathbb{N}^n$ ,  $n \geq 2$
- $\mathcal{X}_{obs}$  – obstacle space,  $\mathcal{X}_{free} = \mathcal{X} \setminus \mathcal{X}_{obs}$  – free space
- $x_{init} \in \mathcal{X}_{free}$  – the initial state,  $x_{goal} \in \mathcal{X}_{free}$  – the goal state
- $\mathcal{X}_{goal} = \left\{ x \in \mathcal{X}_{obs} \mid \|x - x_{goal}\| < r \right\}$  – the goal region
- $\Sigma$  – the set of all feasible paths
- $c(\sigma)$  – the cost function,  $\sigma \in \Sigma$ ,

$$Cost(x_i, x_j) = \|x_i - x_j\|, \quad x_i, x_j \in \mathcal{X}_{free}$$

**Find:** feasible path  $\sigma^* : [0, 1] \rightarrow \mathcal{X}_{free}$

$$\sigma^* = \arg \min_{\sigma \in \Sigma} c(\sigma), \quad s.t. \sigma(0) = x_{init}, \sigma(1) \in \mathcal{X}_{goal}$$

---

\*Project code is available here <https://github.com/akanametov/pathgan>

## 2 Approach

### 2.1 Background

- Sampling-based algorithms solve path planning problems through constructing space-filling trees to search a path  $\sigma$ .
- The tree is built incrementally with samples drawn randomly from the free space  $\mathcal{X}_{free}$
- Drawbacks: the quality of initial solution, the convergence speed

### 2.2 Idea

- Use generative adversarial network (GAN) to learn promising regions and construct heuristic non-uniform sampling distribution  $\mathcal{X}_H \subset \mathcal{X}_{free}$  to reduce sampling space
- Use this heuristic in sampling-base algorithm (e.g., RRT\*)

### 2.3 Algorithm

---

**Algorithm 1:** GAN-based heuristic RRT\*

---

**Input** :  $x_{init}, x_{goal}$ , Map – state space in the form of RGB image  
**Output:**  $G(V, E)$

```
1  $V = x_{init}, E = \emptyset$  ;
2  $\mathcal{S} \leftarrow \text{ROIGenerator}(x_{init}, x_{goal}, \text{Map})$  ;
3  $\mathcal{X}_H \leftarrow \text{Discretization}(\mathcal{S})$ ;
4 for  $i = 1 \dots \text{MAX\_ITER}$  do
5   if  $\text{UseHeuristic} = \text{True}$  then
6     if  $\text{Rand}() > \mu$  then
7        $x_{rand} \leftarrow \text{Non-UniformSample}(\mathcal{X}_H)$ ;
8     else
9        $x_{rand} \leftarrow \text{UniformSample}()$ ;
10    end
11  else
12     $x_{rand} \leftarrow \text{UniformSample}()$ ;
13  end
14 end
15  $x_{nearest} \leftarrow \text{Nearest}(G, x_{rand})$ ;
16  $x_{new} \leftarrow \text{Steer}(x_{nearest}, x_{rand}, \min(\text{MAX\_EDGE\_LEN}, \text{Distance}(x_{nearest}, x_{rand})))$ ;
17 if  $\text{ObstacleFree}(x_{nearest}, x_{rand})$  then
18    $\text{Extend}(G, x_{new})$ ;
19    $\text{Rewire}()$ ;
20   if  $x_{new} \in \mathcal{X}_{goal}$  then
21     Return  $G(V, E)$ ;
22   end
23 end
24 Return Failure;
```

---

ROIGenerator is trained GAN model which outputs image with promising regions (or regions of interest, ROI). Then discretization is performed: coordinates of points from ROI are extracted. Algorithm makes non-uniform sampling from these points with probability  $1 - \mu$  and uniform sampling from all map with probability  $\mu$ .

### 3 GAN-based promising region generation

Trained with large amount of empirical promising region data, the GAN model is able to generate ROI. The input of the model is an RGB image representing the state space Map, start state  $x_{init}$ , and goal state  $x_{goal}$ . The output of the model is also an RGB image where the promising regions are highlighted.

#### 3.1 Dataset generation

---

**Algorithm 2:** Dataset generation

---

**output:** Maps, points, promising regions

- 1 Create some initial maps
- 2 Randomly augment environment maps,  $64 \times 64$
- 3     Color  $\mathcal{X}_{obs}$  into black
- 4     Color  $\mathcal{X}_{free}$  into white
- 5 Randomly choose 100  $x_{init}$  and  $x_{goal}$  states from  $\mathcal{X}_{free}$
- 6     Color  $x_{init}$  into red
- 7     Color  $x_{goal}$  into blue
- 8 Generate promising regions as ground truth
- 9     Randomly select  $x_{init}$ ,  $x_{goal}$  and run RRT 50 times

---

#### 3.2 GAN training and evaluation

---

**Algorithm 3:** GAN training and evaluation

---

**InputGen** :  $\mathcal{P}$  – point images (with  $x_{init}$  and  $x_{goal}$ ),  $\mathcal{M}$  – map images,  $\mathcal{Z}$  – noise data  
**InputDiscr** :  $\mathcal{S}_{gt}$ ,  $\mathcal{S}_p$  – ground truths and predicted Regions of Interest  
**OutputGen** :  $\mathcal{S}_p$  – generated Region of Interest  
**OutputDiscr**: *True* or *Fake*

- 1 **Training;**
- 2 **while** *Generator does not converge* **do**
- 3      $\mathcal{S}_p \leftarrow \text{Generator}(\mathcal{P}, \mathcal{M}, \mathcal{Z})$  ;
- 4      $P_{real\ point} \leftarrow \text{PointDiscriminator}(\mathcal{S}_{gt}, \mathcal{P})$  ;
- 5      $P_{fake\ point} \leftarrow \text{PointDiscriminator}(\mathcal{S}_p, \mathcal{P})$  ;
- 6      $P_{real\ map} \leftarrow \text{MapDiscriminator}(\mathcal{S}_{gt}, \mathcal{M})$  ;
- 7      $P_{fake\ map} \leftarrow \text{MapDiscriminator}(\mathcal{S}_p, \mathcal{M})$  ;
- 8     obtain losses  $\mathcal{L}_{D_{point}}(P_{real\ point}, P_{fake\ point})$ ,  $\mathcal{L}_{D_{map}}(P_{real\ map}, P_{fake\ map})$ ,  $\mathcal{L}_G(\mathcal{S}_{gt}, \mathcal{S}_p)$ ;
- 9     *backprop*
- 10 **end**
- 11 **Evaluation;**
- 12  $\mathcal{S}_p \leftarrow \text{Generator}(\mathcal{P}, \mathcal{M})$  ;
- 13 **Return**  $\mathcal{S}_p$

---

## 4 Experiment results

### 4.1 GAN results

Two GANs were trained on ROIs obtained from RRT pathfinding algorithm.

The first GAN is from the paper and has in five times more parameters compared to the second one - Pix2Pix (ours). Due to this parameters of training like *learning rate* for Generator and Discriminator(s) are slightly different for them:

1. Original GAN (from paper): Generator *learning rate* = 0.0001, Map Discriminator *learning rate* = 0.00005 and Point Discriminator *learning rate* = 0.00005
2. Pix2Pix GAN (ours): Generator *learning rate* = 0.001, Discriminator *learning rate* = 0.0007.

Each of the GANs was trained on 10 *epochs*. Some metrics on test set from generated dataset are shown below:

GAN	mIoU	mDICE	mFID	mIS	number of parameters
Original	<b>70.2%</b>	<b>82.0%</b>	<b>79.7</b>	<b>1.019</b>	21,231,827
Pix2Pix	58.1%	72.2%	91.2	1.017	<b>4,170,477</b>

To check generalization ability of trained GANS, they were tested on unseen MovingAI maps. Results of these ROIs are presented below:

GAN	mIoU	mDICE	mFID	mIS	number of parameters
Original	<b>38.4%</b>	<b>53.8%</b>	<b>88.1</b>	<b>1.014</b>	21,231,827
Pix2Pix	30.8%	46.3%	100.1	1.012	<b>4,170,477</b>

### 4.2 ROI quality evaluation

As proposed in paper we evaluated *connectivity* and *generalization ability*.

**Connectivity** We set the generated ROI for test images as free space and other areas as obstacle space. Then we execute RRT algorithm: if feasible paths inside free space can be found, it means the ROI is connected

**Generalization ability** We resized some maps from Moving AI to  $64 \times 64$ , fed it to trained GANs to generate ROI and evaluated connectivity as above

We measured success rate, % (found connected regions by total number of test maps).

GAN	Generated	MovingAI
Original	<b>65.8%</b>	54.5%
Pix2Pix	65.4%	<b>67.4%</b>

### 4.3 Optimality evaluation

To illustrate the improvement on RRT\*:

1. We randomly choose task for each type of map in the test set
2. We ran RRT\* with uniform sampling and RRT\* with heuristic for 50 times

For first and best found paths the following metrics were collected:

1. Path cost (as sum of euclidean distances between adjacent nodes), time spent in iterations and seconds
2. Nodes added in graph
3. Number of unique nodes sampled (not necessarily added to the graph if obstacles were found inside the new edge)

Also we computed average statistics:

1. Of the all found paths lengths (#nodes) and their costs
2. Of the nodes added in graph and nodes sampled after each 10 iterations

All collected statistics are available by link <https://akanametov.github.io/pathgan/results/>. From these plots we made the following conclusions:

- RRT\* with ROI found first path faster then RRT\* with uniform sampling for both trained models, and difference in time becomes more significant for best path (on average it takes half as many iterations for RRT\* with ROI)
- RRT\* with ROI requires less sampling for finding optimal path
- Paths found by RRT\* with ROI have lower costs and variance
- Although connectivity and CV metrics for unseen MovingAI maps are worse than for our initial dataset, RRT\* with heuristic performs better in most cases

From the line plots we can also observe that RRT\* with heuristic samples less nodes and converges faster to the optimal path. But it can be seen that there is no significant difference between number of nodes in paths found by RRT\* and RRT\* with heuristic.